

Tranalyzer2

Version 0.8.4, Beta, Tarantula



Flow based forensic and network troubleshooting traffic analyzer



Tranalyzer Development Team

Contents

| | | |
|----------|--------------------------------------|----------|
| 1 | Introduction | 1 |
| 1.1 | Getting Tranalyzer | 1 |
| 1.2 | Dependencies | 1 |
| 1.3 | Compilation | 1 |
| 1.4 | Installation | 2 |
| 1.5 | Getting Started | 3 |
| 1.6 | Getting Help | 3 |
| 2 | Tranalyzer2 | 4 |
| 2.1 | Supported Link-Layer Header Types | 4 |
| 2.2 | Enabling/Disabling Plugins | 4 |
| 2.3 | Man Page | 6 |
| 2.4 | Invoking Tranalyzer | 6 |
| 2.5 | ioBuffer.h | 12 |
| 2.6 | main.h | 12 |
| 2.7 | networkHeaders.h | 13 |
| 2.8 | packetCapture.h | 14 |
| 2.9 | tranalyzer.h | 15 |
| 2.10 | bin2txt.h | 21 |
| 2.11 | outputBuffer.h | 21 |
| 2.12 | Tranalyzer2 Output | 22 |
| 2.13 | Final Report | 22 |
| 2.14 | Monitoring Modes During Runtime | 25 |
| 2.15 | Cancellation of the Sniffing Process | 30 |

1 Introduction

Tranalyzer2 is a lightweight flow generator and packet analyzer designed for simplicity, performance and scalability. The program is written in C and built upon the *libpcap* library. It provides functionality to pre- and post-process IPv4/IPv6 data into flows and enables a trained user to see anomalies and network defects even in very large datasets. It supports analysis with special bit coded fields and generates statistics from key parameters of IPv4/IPv6 Tcpdump traces either being live-captured from an Ethernet interface or one or several pcap files. The quantity of binary and text based output of Tranalyzer2 depends on enabled modules, herein denoted as **plugins**. Hence, users have the possibility to tailor the output according to their needs and developers can develop additional plugins independent of the functionality of other plugins.

1.1 Getting Tranalyzer

Tranalyzer can be downloaded from: <https://tranalyzer.com/downloads.html>

1.2 Dependencies

Tranalyzer2 requires **automake**, **libpcap** and **libtool**:

Kali/Ubuntu: `sudo apt-get install automake libpcap-dev libtool make zlib1g-dev`

Arch: `sudo pacman -S automake libpcap libtool zlib`

Fedora/Red Hat/CentOS: `sudo yum install automake libpcap libpcap-devel libtool zlib-devel bzip2`

Gentoo: `sudo emerge autoconf automake libpcap libtool zlib`

OpenSUSE: `sudo zypper install automake gcc libpcap-devel libtool zlib-devel`

Mac OS X: `brew install autoconf automake libpcap libtool zlib`¹

1.3 Compilation

To build Tranalyzer2 and the plugins, run one of the following commands:

- Tranalyzer2 only:
`cd "$T2HOME"; ./autogen.sh tranalyzer2`
(alternative: `cd "$T2HOME/tranalyzer2"; ./autogen.sh`)
- A specific plugin only, e.g., myPlugin:
`cd "$T2HOME"; ./autogen.sh myPlugin`
(alternative 1: `cd "$T2PLHOME/myPlugin"; ./autogen.sh`)
(alternative 2: `cd "$T2HOME/plugins/myPlugin"; ./autogen.sh`)
- Tranalyzer2 and a default set of plugins:
`cd "$T2HOME"; ./autogen.sh`

¹Brew is a packet manager for Mac OS X that can be found here: <https://brew.sh>

- Tranalyzer2 and all the plugins in T2HOME:
`cd "$T2HOME"; ./autogen.sh -a`
- Tranalyzer2 and a custom set of plugins (listed in `plugins.build`) (Section 1.3.1):
`cd "$T2HOME"; ./autogen.sh -b`

where T2HOME points to the root folder of Tranalyzer, i.e., where the file `README.md` is located.

For finer control of which plugins to load, refer to Section 2.2.

Note that if `t2_aliases` is installed, the `t2build` command can be used instead of `autogen.sh`. The command can be run from anywhere, so just replace the above commands with `t2build tranalyzer2`, `t2build myPlugin`, `t2build -a` and `t2build -b`. Run `t2build --help` for the full list of options accepted by the script.

1.3.1 Custom Build

The `-b` option of the `autogen.sh` script takes an optional file name as argument. If none is provided, then the default `plugins.build` is used. The format of the file is as follows:

- Empty lines and lines starting with a ``#'` are ignored (can be used to prevent a plugin from being built)
- One plugin name per row
- Example:

```
# Do not build the tcpStates plugin
#tcpStates

# Build the txtSink plugin
txtSink
```

A `plugins.ignore` file can also be used to prevent specific plugins from being built. A different filename can be used with the `-I` option.

1.4 Installation

The `-i` option of the `autogen.sh` script installs Tranalyzer in `/usr/local/bin` (as `tranalyzer`) and the man page in `/usr/local/man/man1`. Note that root rights are required for the installation.

Alternatively, use the file `t2_aliases` or add the following alias to your `~/.bash_aliases`:

```
alias tranalyzer="$T2HOME/tranalyzer2/src/tranalyzer"
```

where T2HOME points to the root folder of Tranalyzer, i.e., where the file `README.md` is located.

The man page can also be installed manually, by calling (as root):

```
mkdir -p /usr/local/man/man1 && gzip -c man/tranalyzer.1 > /usr/local/man/man1/tranalyzer.1.gz
```

1.4.1 Aliases

The file `t2_aliases` documented in [\\$T2HOME/scripts/doc/scripts.pdf](#) contains a set of aliases and functions to facilitate working with Tranalyzer. To install it, append the following code to `~/.bashrc` or `~/.bash_aliases` (make sure to replace `$T2HOME` with the actual path, e.g., `$HOME/tranalyzer2-0.8.4`):

```
if [ -f "$T2HOME/scripts/t2_aliases" ]; then
    . "$T2HOME/scripts/t2_aliases"          # Note the leading `.'
fi
```

1.5 Getting Started

Run Tranalyzer as follows:

```
tranalyzer -r file.pcap -w outfolder/outprefix
```

For a full list of options, use Tranalyzer `-h` or `--help` option: `tranalyzer -h` or `tranalyzer --help` or refer to the complete documentation.

1.6 Getting Help

1.6.1 Documentation

Tranalyzer and every plugin come with their own documentation, which can be found in the `doc` subfolder. The complete documentation of Tranalyzer2 and all the locally available plugins can be generated by running `make` in `$T2HOME/doc`. The file `t2_aliases` provides the function `t2doc` to allow easy access to the different parts of the documentation from anywhere.

1.6.2 Man Page

If the man page was installed (Section 1.4), then accessing the man page is as simple as calling

```
man tranalyzer
```

If it was not installed, then the man page can be invoked by calling

```
man $T2HOME/tranalyzer2/man/tranalyzer.1
```

1.6.3 Help

For a full list of options, use Tranalyzer `-h` option: `tranalyzer -h`

1.6.4 FAQ

Refer to the complete documentation in `$T2HOME/doc` for a list of frequently asked questions.

1.6.5 Contact

Any feedback, feature requests and questions are welcome and can be sent to the development team via email at:

tranalyzer@rdit.ch

2 Tranalyzer2

Tranalyzer2 is designed in a modular way. Thus, the packet flow aggregation and the flow statistics are separated. While the main program performs the header dissection and flow organisation, the plugins produce specialized output such as packet statistics, mathematical transformations, signal analysis and result file generation.

2.1 Supported Link-Layer Header Types

Tranalyzer handles most PCAP link-layer header types automatically. Some specific types can be analyzed by switching on flags in `linktypes.h`. The following table summarises the link-layer header types handled by Tranalyzer:

| Linktype | Description | Flags |
|-----------------------|--|--------------------|
| DLT_C_HDLC | Cisco PPP with HDLC framing | |
| DLT_C_HDLC_WITH_DIR | Cisco PPP with HDLC framing preceded by one byte direction | |
| DLT_EN10MB | IEEE 802.3 Ethernet (10Mb, 100Mb, 1000Mb and up) | |
| DLT_FRELAY | Frame Relay | |
| DLT_FRELAY_WITH_DIR | Frame Relay preceded by one byte direction | |
| DLT_IEEE802_11 | IEEE802.11 wireless LAN | |
| DLT_IEEE802_11_RADIO | Radiotap link-layer information followed by an 802.11 header | |
| DLT_IPV4 | Raw IPv4 | |
| DLT_IPV6 | Raw IPv6 | |
| DLT_JUNIPER_ATM1 | Juniper ATM1 PIC (experimental) | LINKTYPE_JUNIPER=1 |
| DLT_JUNIPER_ETHER | Juniper Ethernet (experimental) | LINKTYPE_JUNIPER=1 |
| DLT_JUNIPER_PPPOE | Juniper PPPoE PIC (experimental) | LINKTYPE_JUNIPER=1 |
| DLT_LINUX_SLL | Linux “cooked” capture encapsulation | |
| DLT_NULL | BSD loopback encapsulation | |
| DLT_PPI | Per-Packet Information | |
| DLT_PPP | Point-to-Point Protocol | |
| DLT_PPP_SERIAL | PPP in HDLC-like framing | |
| DLT_PPP_WITH_DIR | PPP preceded by one byte direction | |
| DLT_PRISM_HEADER | Prism monitor mode information followed by an 802.11 header | |
| DLT_RAW | Raw IP | |
| DLT_SYMANTEC_FIREWALL | Symantec Enterprise Firewall | |

2.2 Enabling/Disabling Plugins

The plugins are stored under `~/tranalyzer/plugins`. This folder can be changed with the `-p` option.

By default, all the plugins found in the plugin folder are loaded. This behaviour can be changed by altering the value of `USE_PLLIST` in `loadPlugins.h:12`. The valid options are

| USE_PLLIST | Description |
|------------|---|
| 0 | load all plugins in the plugin folder (default) |
| 1 | use a whitelist (loading list) |
| 2 | use a blacklist |

The following sections discuss the various ways to selectively enable/disable plugins.

2.2.1 Default

By default, all the files in the plugin folder named according to the following pattern are loaded:

```
^[0-9]{3}_[a-zA-Z0-9]+.so$
```

To disable a plugin, it must be removed from the plugin folder. A subfolder, e.g., *disabled*, can be used to store unused plugins.

2.2.2 Whitelisting Plugins

If `USE_PLLIST=1`, the whitelist (loading list) is searched under the plugins folder with the name `plugins.txt`. The name can be changed by adapting the value `PL_LIST` in *loadPlugins.h:13*. If the file is stored somewhere else, `Tranalyzer2 -b` option can be used.

The format of the whitelist is as follows (empty lines and lines starting with a '#' are ignored):

```
# This is a comment

# This plugin is whitelisted (will be loaded)
001_protoStats.so

# This plugin is NOT whitelisted (will NOT be loaded)
#010_basicFlow.so
```

Note that if a plugin is not present in the list, it will NOT be loaded.

2.2.3 Blacklisting Plugins

If `USE_PLLIST=2`, the blacklist is searched under the plugins folder with the name `plugins.txt`. The name can be changed by adapting the value `PL_LIST` in *loadPlugins.h:13*. If the file is stored somewhere else, `Tranalyzer2 -b` option can be used.

The format of the blacklist is as follows (empty lines and lines starting with a '#' are ignored):

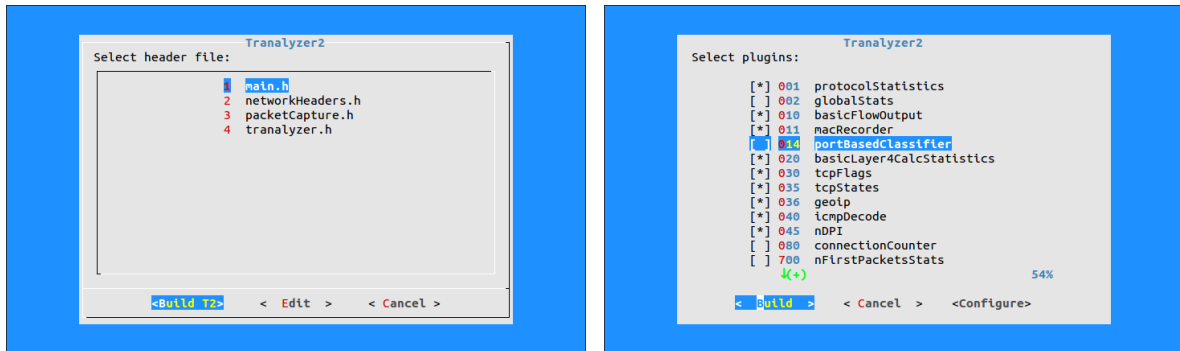
```
# This is a comment

# This plugin is blacklisted (will NOT be loaded)
001_protoStats.so

# This plugin is NOT blacklisted (will be loaded)
#010_basicFlow.so
```

2.2.4 Graphical Configuration and Building of T2 and Plugins

Tranalyzer2 comes with a script named `t2conf` allowing easy configuration of all the plugins through a command line based graphical menu:



Use the arrows on your keyboard to navigate up and down and between the buttons. The first window is only displayed if the `-t2` option is used. The `Edit` and `Configure` buttons will launch a text editor (`$EDITOR` or `vim`² if the environment variable is not defined). The second window can be used to activate and deactivate plugins (toggle the active/inactive state with the space key).

To access the script from anywhere, use the provided `install.sh` script, install `t2_aliases` or manually add the following alias to `~/.bash_aliases`:

```
alias t2conf="$T2HOME/scripts/t2conf/t2conf"
```

Where `$T2HOME` is the folder containing the source code of Tranalyzer2 and its plugins.

A man page for `t2conf` is also provided and can be installed with the `install.sh` script.

2.3 Man Page

If the man page was installed (Section 1.4), then accessing the man page is as simple as calling

```
man tranalyzer
```

If it was not installed, then the man page can be invoked by calling

```
man $T2HOME/tranalyzer2/man/tranalyzer.1
```

2.4 Invoking Tranalyzer

As stated earlier Tranalyzer2 either operates on Ethernet/DAG interfaces or pcap files. It may be invoked using a BPF if only certain flows are interesting. The required arguments are listed below. Note that the `-i`, `-r`, `-R` and `-D` options cannot be used at the same time.

2.4.1 Help

For a full list of options, use the `-h` option: `tranalyzer -h`

Tranalyzer 0.8.2 (Anteater), Tarantula - High performance flow based network traffic analyzer

Usage:

²The default editor can be changed by editing the variable `DEFAULT_EDITOR` (line 7)


```
tranalyzer [OPTION...] <INPUT>
```

Input:

```
-i IFACE      Listen on interface IFACE
-r PCAP       Read packets from PCAP file or from stdin if PCAP is "-"
-R FILE       Process every PCAP file listed in FILE
-D EXPR[:SCHR][,STOP]
               Process every PCAP file whose name matches EXPR, up to an
               optional last index STOP. If STOP is omitted, then Tranalyzer
               never stops. EXPR can be a filename, e.g., file.pcap0, or an
               expression, such as "dump*.pcap00", where the star matches
               anything (note the quotes to prevent the shell from
               interpreting the expression). SCHR can be used to specify the
               the last character before the index (default: 'p')
```

Output:

```
-w PREFIX     Append PREFIX to any output file produced. If omitted, then
               output is diverted to stdout
-W PREFIX[:SIZE][,START]
               Like -w, but fragment flow files according to SIZE, producing
               files starting with index START. SIZE can be specified in bytes
               (default), KB ('K'), MB ('M') or GB ('G'). Scientific notation,
               i.e., 1e5 or 1E5 (=100000), can be used as well. If a 'f' is
               appended, e.g., 10Kf, then SIZE denotes the number of flows.
-l           Print end report in PREFIX_log.txt instead of stdout
-s           Packet forensics mode
```

Optional arguments:

```
-p PATH       Load plugins from path PATH instead of ~/.tranalyzer/plugins
-b FILE       Use plugin list FILE instead of plugin_folder/plugins.txt
-e FILE       Creates a PCAP file by extracting all packets belonging to
               flow indexes listed in FILE
-f FACTOR     Sets hash multiplication factor
-x ID         Sensor ID
-c CPU        Bind tranalyzer to one core. If CPU is 0 then OS selects the
               core to bind
-F FILE       Read BPF filter from FILE

-v           Show the version of the program and exit

-h           Show help options and exit
```

Remaining arguments:

```
BPF          Berkeley Packet Filter command, as in tcpdump
```

2.4.2 -i INTERFACE

Capture data from an Ethernet interface `INTERFACE` (requires *root* privileges). If high volume of traffic is expected, then enable internal buffering in [ioBuffer.h](#).

```
tranalyzer -i eth0 -w out
```

2.4.3 -r FILE

Capture data from a pcap file `FILE`.

```
tranalyzer -r file.pcap -w out
```

The special file `-` can be used to read data from *stdin*. This can be used, e.g., to process compressed pcap files, e.g., *file.pcap.gz*, using the following command:

```
zcat file.pcap.gz | tranalyzer -r - -w out
```

2.4.4 -R FILE

Process all the pcap files listed in `FILE`. All files are being treated as one large file. The life time of a flow can extend over many files. The processing order is defined by the location of the filenames in the text file. The absolute path has to be specified. The `gpl` script documented in `$T2HOME/scripts/scripts.pdf` can be used to generate such a list. All lines starting with a `#` are considered as comments and thus ignored.

```
cd ~/pcap/
$T2HOME/scripts/gpl > pcap_list.txt
tranalyzer -R pcap_list.txt -w out
```

2.4.5 -D FILE[*][.ext]#1[:SCHR][, #2]

Process files in a directory using file start and stop index, defined by `#1` and `#2` respectively. `ext` can be anything, e.g., `.pcap`, and can be omitted. If `#2` is omitted and not in round robin mode, then Tranalyzer2 never stops and waits until the next file in the increment is available. If leading zeroes are used, `#2` defaults to $10^{\text{number_length}} - 1$. Note that only the last occurrence of `SCHR` is considered, e.g., if `SCHR='p'`, then `out.pcap001` will work, but `out001pcap`, will not. with the `:[SCHR]` option a new separation character can be set, superseding `SCHR` defined in [tranalyzer.h](#).

The following variables in [tranalyzer.h](#) can be used to configure this mode:

| Name | Default | Description |
|----------|---------|--|
| RROP | 0 | Activate round robin operations WARNING: if set to 1, then findexer will not work anymore |
| POLLTM | 5 | Poll timing (in seconds) for files |
| MFPTMOUT | 0 | > 0: timeout for poll > POLLTM, 0: no poll timeout |
| SCHR | 'p' | Separating character for file number |

For example, when using `tcpdump` to capture traffic from an interface (`eth0`) and produce 100MB files as follows:

```
sudo tcpdump -C 100 -i eth0 -w out.pcap
```

The following files are generated: *out.pcap*, *out.pcap1*, *out.pcap2*, ..., *out.pcap10*, ...

Then `SCHR` must be set to `'p'`, i.e., the last character before the file number (`out.pcapNUM`) and Tranalyzer must be run as follows:

```
tranalyzer -D out.pcap -w out
```

Or to process files 10 to 100:

```
tranalyzer -D out.pcap10,100 -w out
```

Or to process files 10 to 100 in another format:

```
tranalyzer -D out10.pcap,100 -w out
```

Or to process files from 0 to $2^{32} - 1$ using regex characters:

```
tranalyzer -D "out*.pcap" -w out
```

The last command can be shortened further, the only requirement being the presence of `SCHR` (the last character before the file number) in the pattern:

```
tranalyzer -D "*p" -w out
```

Note the quotes (") which are necessary to avoid preemptive interpretation of regex characters and `SCHR` which **MUST** appear in the pattern. The same configuration can be used for filenames using one or more leading zeros, e.g., *out.pcap000*, *out.pcap001*, *out.pcap002*, ..., *out.pcap010*, ...

The following table summarises the supported naming patterns and the configuration required:

| Filenames | SCHR | Command |
|--|------------------|---|
| <i>out.pcap</i> , <i>out.pcap1</i> , <i>out.pcap2</i> , ... | <code>'p'</code> | <code>tranalyzer -D out.pcap -w out</code> |
| <i>out.pcap00</i> , <i>out.pcap01</i> , <i>out.pcap02</i> , ... | <code>'p'</code> | <code>tranalyzer -D out.pcap00 -w out</code> |
| <i>out0.pcap</i> , <i>out1.pcap</i> , <i>out2.pcap</i> , ... | <code>'t'</code> | <code>tranalyzer -D out0.pcap -w out</code> |
| <i>out00.pcap</i> , <i>out01.pcap</i> , <i>out02.pcap</i> , ... | <code>'t'</code> | <code>tranalyzer -D out00.pcap -w out</code> |
| <i>out_24.04.2016.20h00.pcap</i> , <i>out_24.04.2016.20h00.pcap1</i> , ... | <code>'p'</code> | <code>tranalyzer -D "out*.pcap" -w out</code> |
| <i>out_24.04.2016.20h00.pcap00</i> , <i>out_24.04.2016.20h00.pcap01</i> , ... | <code>'p'</code> | <code>tranalyzer -D "out*.pcap00" -w out</code> |
| <i>out0.pcap</i> , <i>out1.pcap</i> , <i>out2.pcap</i> , ... | <code>'t'</code> | <code>tranalyzer -D out0.pcap:t -w out</code> |
| <i>out.pcap00</i> , <i>out.pcap01</i> , <i>out.pcap02</i> , ... | <code>'p'</code> | <code>tranalyzer -D out.pcap00:p -w out</code> |

2.4.6 -w PREFIX

Use a `PREFIX` for all output file types. The number of files being produced vary with the number of activated plugins. The file suffixes are defined in the file `tranalyzer.h` (see Section 2.9.13) or in the header files for the plugins. If you forget to specify an output file, Tranalyzer will use the input interface name or the file name as file prefix and print the flows to *stdout*. Thus, Tranalyzer output can be piped into other command line tools, such as netcat in order to produce centralized logging to another host or an AWK script for further post processing without intermediate writing to a slow disk storage.

2.4.7 -W PREFIX[:SIZE][,START]

This option allows the fragmentation of flow files produced by Tranalyzer independent of the input mode. The expression before the ':' is the output prefix, the expression after the ':' denotes the maximal file size for each fragment and the number after the ',' denotes the start index of the first file. If omitted it defaults to 0. The size of the files can be specified in bytes (default), KB ('K'), MB ('M') or GB ('G'). Scientific notation, i.e., 1e5 or 1E5 (=100000), can be used as well. If no size is specified, the default value of 500MB, defined by OFRWFILELN in [tranalyzer.h](#) is used. If no size is specified, then the ':' can be omitted. The same happens if no start index is specified. If an additional 'f' is appended the unit is flow count. This enables the user to produce file chunks containing the same amount of flows. Some typical examples are shown below.

| Command | Fragment Size | Start Index | Output Files |
|--|---------------|-------------|-------------------|
| tranalyzer -r nudel.pcap -W out:1.5E9,10 | 1.5GB | 10 | out10, out11, ... |
| tranalyzer -r nudel.pcap -W out:1.5e9,5 | 1.5GB | 5 | out5, out6, ... |
| tranalyzer -r nudel.pcap -W out:1.5G,1 | 1.5GB | 1 | out1, out2, ... |
| tranalyzer -r nudel.pcap -W out:5000K | 0.5MB | 0 | out0, out1, ... |
| tranalyzer -r nudel.pcap -W out:5Kf | 5000 Flows | 0 | out0, out1, ... |
| tranalyzer -r nudel.pcap -W out:180M | 180MB | 0 | out0, out1, ... |
| tranalyzer -r nudel.pcap -W out:2.5G | 2.5GB | 0 | out0, out1, ... |
| tranalyzer -r nudel.pcap -W out,5 | OFRWFILELN | 0 | out0, out1, ... |
| tranalyzer -r nudel.pcap -W out | OFRWFILELN | 0 | out0, out1, ... |

2.4.8 -l

All Tranalyzer command line and report output is diverted to the log file: PREFIX_log.txt. Fatal error messages still appear on the command line.

2.4.9 -s

Initiates the packet mode, where a file with the suffix PREFIX_packets.txt is created. The content of the file depends on the plugins loaded. The display of the packet number (first column is controlled by SPKTM_PACKETNO in [main.h](#). The layer 7 payload can be displayed in hexadecimal and/or as characters by using the SPKTM_PCNTX and SPKTM_PCNTC respectively. A tab separated header description line is printed at the beginning of the packet file. The first two lines then read as follows:

| %pktNo | time | pktIAT | duration | flowInd | flowStat | numHdrDesc | hdrDesc | ethVlanID | | | | | | | | |
|--------|-------------------|--|-------------------|-------------------|--------------|-------------|------------------------|-----------|---------|--------|------|------|------|------|------|------|
| | ethType | srcMac | dstMac | srcIP4 | srcPort | dstIP4 | dstPort | l4Proto | ipTOS | | | | | | | |
| | ipID | ipIDDiff | ipFrag | ipTTL | ipHdrChkSum | ipCalChkSum | l4HdrChkSum | | | | | | | | | |
| | l4CalChkSum | ipFlags | pktLen | ipOptLen | ipOpts | seq | ack | seqDiff | ackDiff | | | | | | | |
| | seqPktLen | ackPktLen | tcpStat | tcpFlags | tcpAnomaly | tcpWin | tcpOptLen | tcpOpts | | | | | | | | |
| | l7Content | | | | | | | | | | | | | | | |
| ... | | | | | | | | | | | | | | | | |
| 25 | 1291753225.446846 | 0.000000 | 0.000000 | 23 | 0x00006000 | 6 | eth:vlan:mpls{2}:ipv4: | | | | | | | | | |
| | tcp 20 | 0x0800 | 00:90:1a:41:fa:45 | 00:13:c4:52:4a:07 | 188.62.56.56 | 62701 | | | | | | | | | | |
| | 212.243.221.241 | 80 | 6 | 0x00 | 0x26f6 | 0 | 0x4000 | 62 | 0x6ca6 | 0x6ca6 | 0 | | | | | |
| | x0247 | 0x0247 | 0x0040 | 460 | 0 | 0xb2a08909 | 0x90314073 | 0 | 0 | 0 | 0 | | | | | |
| | x59 | 0x18 | 0x0000 | 65535 | 12 | 0x01 | 0x01 | 0x08 | 0x0a | 0x29 | 0x2d | 0xc3 | 0x6e | 0x83 | 0x63 | 0xc5 |
| | 0x76 | GET /images/I/01TnJ0+mhnL.png HTTP/1.1\r\nHost: ecx.images-amazon.com\r\nUser-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10.6; de; rv:1.9.2.8) Gecko/20100722 Firefox/3.6.8\r\nAccept: image/png, image/*; q=0.8, */*; q=0.5\r\nAccept-Language: de-de, de; q=0.8, en-us; q=0.5, en; | | | | | | | | | | | | | | |

```
q=0.3\r\nAccept-Encoding: gzip, deflate\r\nAccept-Charset: ISO-8859-1, utf-8;q=0.7,*;q=0.7\r\n
nKeep-Alive: 115\r\nConnection: keep-alive\r\nReferer: http://z-ecx.images-amazon.com/images/I
/11J5cf408UL.css\r\n\r\n
...
```

2.4.10 -p FOLDER

Changes the plugin folder from standard `~/.tranalyzer/plugins` to `FOLDER`.

2.4.11 -b FILE

Changes the plugin blacklist file from `plugin_folder/plugin_blacklist.txt` to `FILE`, where `plugin_folder` is either `~/.tranalyzer/plugins` or the folder specified with the `-p` option.

2.4.12 -e FLOWINDEXFILE

Denotes the filename and path of the flow index file when the `pcapd` plugin is loaded. The path and name of the pcap file depends on `FLOWINDEXFILE`. If omitted the default names for the PCAP file are defined in `pcapd.h`. The format of the `FLOWINDEXFILE` is a list of 64 bit flow indices which define the packets to be extracted from the pcap being read by the `-r` option. In general the user should use a plain file with the format displayed below:

```
# Comments (ignored)
% Flow file info (ignored)
30
3467
656697
5596
```

For more information on the `pcapd` plugin please refer to its [documentation](#).

2.4.13 -f HASHFACTOR

Sets and superseeds the `HASHFACTOR` constant in [tranalyzer.h](#).

2.4.14 -x SENSORID

Each T2 can have a separate sensor ID which can be listed in a flow file in order to differentiate flows originating from several interfaces during post processing, e.g., in a DB. If not specified `T2_SENSORID` (666), defined in [tranalyzer.h](#), will be the default value.

2.4.15 -c CPU

Bind Tranalyzer to core number `CPU`; if `CPU == 0` then the operating system selects the core to bind.

2.4.16 -F FILE

Read BPF filter from `FILE`. A filter can span multiple lines and can be commented using the ``#'` character (everything following a ``#'` is ignored).

2.4.17 BPF Filter

A Berkeley Packet Filter (BPF) can be specified at any time in order to reduce the amount of flows being produced and to increase speed during life capture ops. All rules of pcap BPF apply.

2.5 ioBuffer.h

| Name | Default | Description |
|---|---------|---|
| ENABLE_IO_BUFFERING | 0 | Enables buffering of the packets in a queue |
| If ENABLE_IO_BUFFERING == 1, the following flags are available: | | |
| IO_BUFFER_FULL_WAIT_MS | 200 | Number of milliseconds to wait if queue is full |
| IO_BUFFER_SIZE | 8192 | Maximum number of packets that can be stored in the buffer (power of 2) |
| IO_BUFFER_MAX_MTU | 2048 | Maximum size of a packet (divisible by 4) |

2.6 main.h

The monitoring mode can be configured with the following constants:

| Name | Default | Description |
|--|-------------|---|
| SPKTM_PACKETNO | 1 | Print the packet number |
| SPKTM_PCNTC | 1 | Print L7 content as characters |
| SPKTM_PCNTX | 0 | Print L7 content as hex |
| MIN_MAX_ESTIMATE | 0 | Min/Max bandwidth statistics |
| The following flags control the monitoring mode: | | |
| MONINTTHR | 1 | Activate threaded interrupt handling |
| MONINTBLK | 0 | Block interrupts in main loop during packet processing, disables MONINTTHR |
| MONINTSYNC | 1 | Synchronized print statistics |
| MONINTTMPCP | 0 | 0: real time base, 1: pcap time base |
| MONINTTMPCP_ON | 0 | Startup monitoring. 1: on 0: off (require MONINTTMPCP=1) |
| MONINTV | 1 | ≥ 1 sec interval of monitoring output if USR2 is sent or MONINTTMPCP=0 |
| MONPROTMD | 1 | 0: report protocol numbers; 1: report protocol names |
| MONPROTFL | "proto.txt" | proto file |

The MONPROTL2 and MONPROTL3 flags can be used to configure the L2 and L3 protocols to monitor. Their default values are

- MONPROTL2: 0x0042, 0x00fe, ETHERTYPE_ARP, ETHERTYPE_RARP, ETHERTYPE_IP, ETHERTYPE_IPV6
- MONPROTL3: L3_ICMP, L3_IGMP, L3_TCP, L3_UDP, L3_GRE, L3_ICMP6, L3_SCTP

2.7 networkHeaders.h

| Name | Default | Description | Flags |
|-----------------|---------|---|------------------|
| IPV6_ACTIVATE | 2 | 0: IPv4 only 1: IPv6 only 2: Dual mode | |
| ETH_ACTIVATE | 1 | 0: No Ethernet flows 1: Activate Ethernet flows 2: Also use Ethernet addresses for IPv4/6 flows | |
| SCTP_ACTIVATE | 0 | SCTP protocol decoder for stream → flow generation is activated | |
| SCTP_STATFINDEX | 0 | 1: finindex constant for all SCTP streams in a packet 0: finindex increments | SCTP_ACTIVATE=1 |
| MULTIPKTSUP | 0 | Multi-packet suppression (discard duplicated packets) | IPV6_ACTIVATE=0 |
| T2_PRI_HDRDESC | 1 | 1: keep track of the headers traversed | |
| T2_HDRDESC_AGGR | 1 | 1: aggregate repetitive headers, e.g., vlan{2} | T2_PRI_HDRDESC=1 |
| T2_HDRDESC_LEN | 128 | max length of the headers description | T2_PRI_HDRDESC=1 |

2.8 *packetCapture.h*

The config file *packetCapture.h* provides control about the packet capture and packet structure process of Tranalyzer2. The most important fields are described below. Please note that after changing any value in define statements a rebuild is required. Note that the `PACKETLENGTH` switch controls the `packetLength` variable in the packet structure, from where the packet length is measured from. So statistical plugins such as `basicStats` can have a layer dependent output. If only L7 length is needed, use the `packetL7length` variable in the packet structure.

| Name | Default | Description | Flags |
|-------------------------------|---------|--|------------------------------|
| <code>PACKETLENGTH</code> | 3 | 0: including L2, L3 and L4 header 1: including L3 and L4 header 2: including L4 header 3: only higher layer payload (Layer 7) | |
| <code>FRGIPPKTLENVIEW</code> | 1 | 0: IP header stays with 2nd++ fragmented packets 1: IP header stripped from 2nd++ fragmented packets | <code>PACKETLENGTH=1</code> |
| <code>NOLAYER2</code> | 0 | 0: Automatic L3 header discovery 1: Manual L3 header positioning | |
| <code>NOL2_L3HDROFFSET</code> | 0 | Offset of L3 header | <code>NOLAYER2=1</code> |
| <code>MAXHDCNT</code> | 5 | Maximal header count (MUST be ≥ 3) | <code>IPV6_ACTIVATE=1</code> |

2.9 tranalyzer.h

| Name | Default | Description | Flags |
|-------------------------|---------|---|--------------|
| REPSUP | 0 | Activate alive mode | |
| PID_FNM_ACT | 0 | Save the PID into a file <code>PID_FNM</code> (default: "tranalyzer.pid") | |
| DEBUG | 0 | 0: no debug output 1: debug output which occurs only once or very seldom 2: + debug output which occurs in special situations, but not regularly 3: + debug output which occurs regularly (every packet) | |
| VERBOSE | 2 | 0: no output 1: basic pcap report 2: + full traffic statistics 3: + info about fragmentation anomalies | |
| MEMORY_DEBUG | 0 | 0: no memory debug 1: detect leaks and overflows (see <i>utils/memdebug.h</i>) | |
| NO_PKTS_DELAY_US | 1000 | If no packets are available, sleep for <i>n</i> microseconds | |
| NON_BLOCKING_MODE | 1 | Non-blocking mode | |
| MAIN_OUTPUT_BUFFER_SIZE | 1000000 | Size of the main output buffer | |
| SNAPLEN | BUFSIZ | Snapshot length (live capture) | |
| CAPTURE_TIMEOUT | 1000 | Read timeout in milliseconds (live capture) | |
| ENABLE_BPF_OPTIMIZATION | 0 | Optimize BPF filters | |
| TSTAMP_PREC | 0 | Timestamp precision: 0: microseconds, 1: nanoseconds | |
| TSTAMP_UTC | 1 | Time representation: 0: localtime, 1: UTC | |
| TSTAMP_R_UTC | 0 | Time report representation: 0: localtime, 1: UTC | |
| ALARM_MODE | 0 | Only output flow if an alarm-based plugin fires | |
| ALARM_AND | 0 | 0: logical OR, 1: logical AND | ALARM_MODE=1 |
| FORCE_MODE | 0 | Parameter induced flow termination (NetFlow mode) | |
| BLOCK_BUF | 0 | Block unnecessary buffer output when non Tranalyzer format event based plugins are active | |

| Name | Default | Description | Flags |
|-----------------|---------|---|------------------|
| PLUGIN_REPORT | 1 | Enable plugins to contribute to Tranalyzer end report | |
| DIFF_REPORT | 0 | 0: absolute Tranalyzer command line USR1 report 1: differential report | |
| MACHINE_REPORT | 0 | USR1 report: 0: human compliant, 1: machine compliant | |
| REPORT_HIST | 0 | Store statistical report history in <code>REPORT_HIST_FILE</code> after shutdown and reload it when restarted | |
| ESOM_DEP | 0 | Allow plugins to globally access plugin dependent variables | |
| AYIYA | 1 | Activate AYIYA processing | |
| GENEVE | 1 | Activate GENEVE processing | |
| TEREDO | 1 | Activate TEREDO processing | |
| L2TP | 1 | Activate L2TP processing | |
| GRE | 1 | Activate GRE processing | |
| GTP | 1 | Activate GTP (GPRS Tunneling Protocol) processing | |
| VXLAN | 1 | Activate VXLAN (Virtual eXtensible Local Area Network) processing | |
| IPIP | 1 | Activate IPv4/6 in IPv4/6 processing | |
| ETHIP | 1 | Activate Ethernet within IP IPv4/6 processing | |
| CAPWAP | 1 | Activate CAPWAP processing | |
| CAPWAP_SWAP_FC | 1 | Swap frame control (required for Cisco) | |
| LWAPP | 1 | Activate LWAPP processing | |
| LWAPP_SWAP_FC | 1 | Activate LWAPP processing (required for Cisco) | |
| FRAGMENTATION | 1 | Activate fragmentation processing | |
| FRAG_HLST_CRFT | 1 | Enables crafted packet processing | FRAGMENTATION=1 |
| FRAG_ERROR_DUMP | 0 | Dumps flawed fragmented packet to <code>stdout</code> | FRAGMENTATION=1 |
| IPVX_INTERPRET | 0 | Interpret bogus IPvX packets | |
| ETH_STAT_MODE | 0 | Whether to use the innermost (0) or outermost (1) layer 2 type for the statistics | |
| RELTIME | 0 | 0: absolute time 1: relative time | |
| FDURLIMIT | 0 | If > 0 , force flow life span to $n \pm 1$ seconds | |
| FLOW_TIMEOUT | 182 | Flow timeout after a packet is not seen after n seconds | |
| HASH_AUTOPILOT | 1 | 0: terminate when main hash map is full 1: flushes oldest <code>NUMFLWRM</code> flow(s) when main hash is full | |
| NUMFLWRM | 1 | Number of flows to flush when main hash map is full | HASH_AUTOPILOT=1 |

2.9.1 -D constants

the following constants influence the file name convention:

| Name | Default | Description |
|--------|---------|--------------------------------------|
| RROP | 0 | round robin operations |
| POLLTM | 5 | poll timing for files |
| SCHR | 'p' | separating character for file number |

2.9.2 alive signal

The alive signal is a derivate of the passive monitoring mode by the USR1 signal, where the report is deactivated. If REPSUP=1 then only the command defined by REPCMDAS/W is sent to the control program defined by ALVPROG as defined below:

| Name | Default | Description |
|----------|---|--|
| REPSUP | 0 | 0: alive mode off, 1: alive mode on, monitoring report suppressed |
| ALVPROG | "t2alive" | name of control program |
| REPCMDAS | "a='pgrep ALVPROG'; \ if [\$a]; then kill -USR1 \$a; fi" | alive and stall USR1 signal (no packets) |
| REPCMDAW | "a='pgrep ALVPROG'; \ if [\$a]; then kill -USR2 \$a; fi" | alive and well USR2 signal (working) |

If T2 crashes or is stopped a syslog message is issued by the t2alive daemon. Same if T2 gets started.

2.9.3 FORCE_MODE

A 1 enables the force mode which enables any plugin to force the output of flows independent of the timeout value. Hence, Cisco NetFlow similar periodic output can be produced or overflows of counters can produce a flow and restart a new one.

2.9.4 ALARM_MODE

A 1 enables the alarm mode which differs from the default flow mode by the plugin based control of the Tranalyzer core flow output. It is useful for classification plugins generating alarms, thus emulating alarm based SW such as Snort, etc. The default value is 0. The plugin sets the global output suppress variable supOut=1 in the *onFlowTerminate()* function before any output is generated. This mode also allows multiple classification plugins producing an 'AND' or an 'OR' operation if many alarm generating plugins are loaded. The variable ALARM_AND controls the logical alarm operation. A sample code which has to be present at the beginning of the *onFlowTerminate()* function is shown below:

```
#if ALARM_MODE == 1
#if ALARM_AND == 0
    if (!Alarm) supOut = 0;
#else // ALARM_AND == 1
    if (!Alarm) {
        supOut = 1;
        return;
    }
#endif // ALARM_AND
#endif // ALARM_MODE == 1
```

Figure 1: A sample code in the *onFlowTerminate()* routine

2.9.5 BLOCK_BUF

if set to '1' unnecessary buffered output from all plugins is blocked when non Tranalyzer format event based plugins are active: e.g. syslog, arcsight and text-based or binary output plugins are not loaded.

2.9.6 Report Modes

Tranalyzer provides a user interrupt based report and a final report. The interrupt based mode can be configured in a variety of ways being defined below.

| Name | Default | Description |
|----------------|---------|--|
| PLUGIN_REPORT | 0 | enable plugins to contribute to the tranalyzer command line end report |
| DIFF_REPORT | 0 | 1: differential, 0: Absolute tranalyzer command line USR1 report |
| MACHINE_REPORT | 0 | USR1 Report 1: machine compliant; 0: human compliant |

The following interrupts are being caught by Tranalyzer2:

| Signal Name | Description |
|-------------|---|
| SIGINT | like ^C terminates new flow production ³ |
| SIGTERM | terminates tranalyzer |
| SIGUSR1 | prints statistics report |
| SIGUSR2 | toggles repetitive statistics report |

2.9.7 State and statistical save mode

T2 is capable to preserve its internal statistical state and certain viable global variables, such as the index.

| Name | Default | Description |
|------------------|-----------------|---|
| REPORT_HIST | 0 | Store statistical report history after shutdown, reload it upon restart |
| REPORT_HIST_FILE | "stat_hist.txt" | default statistical report history filename |

The history file is stored by default under `./tranalyzer/plugins` or under the directory defined by a `-p` option.

2.9.8 L2TP

A '1' activates the L2TP processing of the Tranalyzer2 core. All L2TP headers either encapsulated in MPLS or not will be processed and followed down via PPP headers to the IP header and then passed to the IP processing. The default value of the variable is '0'. Then the stack will be parsed until the first IP header is detected. So all L2TP UDP headers having src and dest port 1701 will be processed as normal UDP packets.

2.9.9 GRE

A '1' activates the L3 General Routing Encapsulation (L4proto=47) processing of the Tranalyzer2 core. All GRE headers either encapsulated in MPLS or not will be processed and followed down via PPP headers to the IP header and then passed to the IP processing. The default value of the variable is 0. Then the stack will be parsed until the first IP header is detected. If the following content is not existing or compressed the flow will contain only L4proto = 47 information.

³If two SIGINT interrupts are being sent in short order Tranalyzer will be terminated instantly.

2.9.10 FRAGMENTATION

A '1' activates the fragmentation processing of the Tranalyzer2 core. All packets following the header packet will be assembled in the same flow. The core and the plugin tcpFlags will provide special flags for fragmentation anomalies. If FRAGMENTATION is set to 0 only the initial fragment will be processed; all later fragments will be ignored.

2.9.11 FRAG_HLST_CRFT

A '1' enables crafted packet processing even when the lead fragment is missing or packets contain senseless flags as being used in attacks or equipment failure.

2.9.12 FRAG_ERROR_DUMP

A '1' activates the dump of packet information on the command line for time based identification of ill-fated or crafted fragments in tcpdump or wireshark. It provides the Unix timestamp, the six tuple, IPID and fragID as outlined in figure below.

| MsgType | msg | time | vlan | srcIP | srcPort | dstIP | dstPort | proto | fragID | fragOffset |
|---------|------------------------------------|-------------------|------|----------------|---------|----------------|---------|-------|--------|------------|
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.449690 | 20 | 92.104.181.154 | 42968 | 93.144.66.3 | | | | |
| | 52027 17 - 0x191F 0x00A0 | | | | | | | | | |
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.482611 | 20 | 92.104.181.154 | 43044 | 93.144.66.3 | | | | |
| | 1719 17 - 0x1922 0x00A0 | | | | | | | | | |
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.492830 | 20 | 92.104.181.154 | 55841 | 93.144.66.3 | | | | |
| | 28463 17 - 0x1923 0x00A0 | | | | | | | | | |
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.503955 | 20 | 92.104.181.154 | 25668 | 93.144.66.3 | | | | |
| | 8137 17 - 0x1924 0x00A0 | | | | | | | | | |
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.551094 | 20 | 92.105.93.227 | 41494 | 24.218.128.232 | | | | |
| | 27796 17 - 0x5A21 0x00A0 | | | | | | | | | |
| [WRN] | packetCapture: 1. frag not found @ | 1291753225.639627 | 20 | 86.51.18.243 | 38824 | 92.105.108.208 | | | | |
| | 55133 17 - 0x0DAE 0x00AC | | | | | | | | | |

Figure 2: A sample report on stdout for packets with an elusive first fragment

WARNING: If FRAG_HLST_CRFT == 1 then every fragmented headerless packet will be reported!

2.9.13 *_SUFFIX

This constant defines the suffix of all plugin output files. For example if you specify the output *foo.foo* (with the *-w* option), the generated file for the per-packet output will be in the default setting *foo.foo_packets*.

2.9.14 RELTIME

RELTIME renders all time based plugin output into relative to the beginning of the pcap or start of packet capture. In *-D* or *-R* read operation the first file defines the start time.

2.9.15 FLOW_TIMEOUT

This constant specifies the default time in seconds (182) after which a flow will be considered as terminated since the last packet is captured. Note: Plugins are able to change the timeout values of a flow. For example the *tcpStates* plugin adjusts the timeout of a flow according to the TCP state machine. A reduction of the flow timeout has an effect on the necessary flow memory defined in *HASHCHAINTABLE_SIZE*, see below.

2.9.16 FDURLIMIT

FDURLIMIT defines the maximum flow duration in seconds which is then forced to be released. It is a special force mode for the duration of flows and a special feature for Dalhousie University. If `FDURLIMIT > 0` then `FLOW_TIMEOUT` is overwritten if `FURLIMIT` seconds are reached.

2.9.17 HASHFACTOR

A factor to be multiplied with the `HASHTABLE_SIZE` and `HASHCHAINEDTABLE_SIZE` described below. It facilitates the correct setting of the hash space. Moreover, if T2 runs out of hash it will give an upper estimate the user can choose for `HASHFACTOR`. Set it to this value, recompile and rerun T2. This constant is superseded by the `-f` option.

2.9.18 HASHTABLE_SIZE

The number of buckets in the hash table. As a separate chaining hashing method is used, this value does not denote the amount of elements the hash table is able to manage! The larger, the less likely are hash collisions. The current default value is 2^{18} . Its value should be selected at least two times larger as the value of `HASHCHAINEDTABLE_SIZE` discussed in the following chapter.

2.9.19 HASHCHAINEDTABLE_SIZE

Specifies the amount of flows the main hash table is able to manage. The default value is 2^{19} , so roughly half the size of `HASHTABLE_SIZE`. T2 supplies information about the hash space in memory in: Max number of IPv4 flows in memory: 113244 (50.220%). Together with the amount of traffic already processed the total value can be computed. An example is given in Figure 1.

2.9.20 HASH_AUTOPILOT

Default 1. Avoids overrun of main hash, flushes oldest flow on every flowInsert if hashmap is full. 0 disables hash overrun protection. If speed is an issue avoid overruns by invoking T2 with a laged `-f` option value, as T2 recommends.

2.9.21 Aggregation Mode

The aggregation mode enables the user to confine certain IP, port or protocol ranges into a single flow. The variable `AGGREGATIONFLAG` in [tranalyzer.h](#) defines a bit field which enables specific aggregation modes according to the six tuple values listed below.

| Aggregation Flag | Value |
|------------------|-------|
| L4PROT | 0x01 |
| DSTPORT | 0x02 |
| SRCPORT | 0x04 |
| DSTIP | 0x08 |
| SRCIP | 0x10 |
| VLANID | 0x20 |
| SUBNET | 0x80 |

If a certain aggregation mode is enabled the following variables in [tranalyzer.h](#) define the aggregation range.

| Aggregation Flag | Type | Description |
|------------------|----------|--------------------------------|
| SRCIP4CMSK | uint8_t | src IPv4 aggregation CIDR mask |
| DSTIP4CMSK | uint8_t | dst IPv4 aggregation CIDR mask |
| SRCIP6CMSK | uint8_t | src IPv6 aggregation CIDR mask |
| DSTIP6CMSK | uint8_t | dst IPv6 aggregation CIDR mask |
| SRCPORTLW | uint16_t | src port lower bound |
| SRCPORTHW | uint16_t | src port upper bound |
| DSTPORTLW | uint16_t | dst port lower bound |
| DSTPORTHW | uint16_t | dst port upper bound |

2.10 bin2txt.h

| Name | Default | Description |
|------------------------|---------|---|
| HEX_CAPITAL | 0 | hex output: 0: lower case; 1: upper case |
| IP4_FORMAT | 0 | IPv4 addresses representation: 0: normal, 1: normalized (padded with zeros), 2: one 32-bits hex number 3: one 32-bits unsigned number |
| IP6_FORMAT | 0 | IPv6 addresses representation: 0: compressed, 1: uncompressed, 2: one 128-bits hex number, 3: two 64-bits hex numbers |
| MAC_FORMAT | 0 | MAC addresses representation: 0: normal (edit MAC_SEP to change the separator), 1: one 64-bits hex number, |
| MAC_SEP | ": " | Separator to use in MAC addresses: 11:22:33:44:55:66 |
| TFS_EXTENDED_HEADER | 0 | Extended header in flow file |
| B2T_TIME_IN_MICRO_SECS | 1 | Time precision: 0: nanosecs, 1: microsecs |
| TFS_NC_TYPE | 1 | Types in header file: 0: numbers, 1: C types |
| TFS_SAN_UTF8 | 1 | Activates the UTF-8 sanitizer for strings |
| B2T_TIMESTR | 0 | Print unix timestamps as human readable dates |
| HDR_CHR | "%" | start character(s) of comments |
| SEP_CHR | "\t" | column separator in the flow file ";", ". ", "_ " and "\" should not be used |
| JSON_KEEP_EMPTY | 0 | Whether or not to output empty fields |
| JSON_PRETTY | 0 | Whether to add spaces to make the output more readable |

2.11 outputBuffer.h

| Name | Default | Description | Flags |
|---------------|---------|--|-------|
| BUF_DATA_SHFT | 0 | Adds for each binary output record the length and shifts the record by n uint32_t words to the right | |

| Name | Default | Description | Flags |
|------------------|---------|---|--------------------|
| OUTBUF_AUTOPILOT | 1 | (see <code>binSink</code> and <code>socketSink</code> plugin) Automatically increase the output buffer when required | |
| OUTBUF_MAXSIZE_F | 5 | Maximal factor to increase the output buffer size to | OUTBUF_AUTOPILOT=1 |

2.12 Tranalyzer2 Output

As stated before, the functionality and output of Tranalyzer2 is defined by the activated plugins. Basically, there are two ways a plugin can generate output. First, it can generate its own output file and write any arbitrary content into any stream. The second way is called standard output or per-flow output. After flow termination Tranalyzer2 provides an output buffer and appends the direction of the flow to it. For example, in case of textual output, an "A" flow is normally followed by a "B" flow or if the "B" flow does not exist it is followed by the next "A" flow. Then, the output buffer is passed to the plugins providing their per-flow output. Finally the buffer is sent to the activated output plugins. This process repeats itself for the "B" flow. For detailed explanation about the functionality of the output plugins refer to the section plugins.

2.12.1 Hierarchical Ordering of Numerical or Text Output

Tranalyzer2 provides a hierarchical ordering of each output. Each plugin controls the:

- volume of its output
- number of values or bins
- hierarchical ordering of the data
- repetition of data substructures

Thus, complex structures such as lists or matrices can be presented in a single line.

The following sample of text output shows the hierarchical ordering for four data outputs, separated by tabulators:

| | | | | |
|---|-----|-------------|---------|---------------------------------|
| A | 0.3 | 2.0_3.4_2.1 | 2;4;2;1 | (1_2_9)_(1_3_1)_(7_5_3)_(2_3_7) |
|---|-----|-------------|---------|---------------------------------|

The A indicates the direction of the flow; in this case it is the initial flow. The next number denotes a singular descriptive statistical result. Output number two consists of three values separated by “_” characters. Output number three consists of one value, that can be repeated, indicated by the character “;”. Output number four is a more complex example: It consists of four values containing three subvalues indicated by the braces. This could be interpreted as a matrix of size 4×3 .

2.13 Final Report

Standard configuration of Tranalyzer2 produces a statistical report to *stdout* about timing, packets, protocol encapsulation type, average bandwidth, dump length, etc. A sample report including some current protocol relevant warnings is depicted in the figure below. Warnings are not fatal hence are listed at the end of the statistical report when Tranalyzer2 terminates naturally. The *Average total Bandwidth* estimation refers to the processed bandwidth during the data acquisition process. It is only equivalent to the actual bandwidth if the total packet length including all encapsulations is not truncated and all traffic is IP. The *Average IP Traffic Bandwidth* is an estimate comprising all IP traffic actually present on the wire. Plugins can report extra information when `PLUGIN_REPORT` is activated. This report can be saved in a file, by using one of the following command:


```

tranalyzer -r file.pcap -w out -l (See Section 2.4.8)
tranalyzer -r file.pcap -w out | tee out_stdout.txt
tranalyzer -r file.pcap -w out > out_stdout.txt

```

Both commands will create a file `out_stdout.txt` containing the report. The only difference between those two commands is that the first one still outputs the report to stdout.

Fatal errors regarding the invocation, configuration and operation of Tranalyzer2 are printed to *stderr* after the plugins are loaded, thus before the processing is activated, see the *Hash table error* example in Listing 1. These errors terminate Tranalyzer2 immediately and are located before the final statistical report as being indicated by the “*Shutting down...*” key phrase. If the final report is to be used in a following script a pipe can be appended and certain lines can be filtered using `grep` or `awk`.

```

$ ./tranalyzer -r ~/data/knoedel.pcap -w ~/results/
=====
Tranalyzer 0.8.2 (Anteater), Tarantula. PID: 16298
=====
[INF] Creating flows for L2, IPv4, IPv6
Active plugins:
  01: protoStats, 0.8.2
  02: basicFlow, 0.8.2
  03: macRecorder, 0.8.2
  04: portClassifier, 0.8.2
  05: basicStats, 0.8.3
  06: tcpFlags, 0.8.2
  07: tcpStates, 0.8.2
  08: icmpDecode, 0.8.2
  09: dnsDecode, 0.8.3
  10: httpSniffer, 0.8.2
  11: connStat, 0.8.2
  12: txtSink, 0.8.2
[INF] basicFlow: IPv4 Ver: 3, Rev: 20190114, Range Mode: 0, subnet ranges loaded: 2821502 (2.82 M)
[INF] basicFlow: IPv6 Ver: 3, Rev: 20190114, Range Mode: 0, subnet ranges loaded: 36123 (36.12 K)
Processing file: /home/user/data/knoedel.pcap
Link layer type: Ethernet [EN10MB/1]
Dump start: 1291753225.446732 sec (Tue 07 Dec 2010 20:20:25 GMT)
[WRN] snapL2Length: 1550 - snapL3Length: 1484 - IP length in header: 1492
[WRN] Hash Autopilot: main HashMap full: flushing 1 oldest flow(s)
[INF] Hash Autopilot: Fix: Invoke Tranalyzer with '-f 5'
Dump stop : 1291753452.373884 sec (Tue 07 Dec 2010 20:24:12 GMT)
Total dump duration: 226.927152 sec (3m 46s)
Finished processing. Elapsed time: 284.541828 sec (4m 44s)
Finished unloading flow memory. Time: 295.292141 sec (4m 55s)
Percentage completed: 100.00%
Number of processed packets: 53982409 (53.98 M)
Number of processed bytes: 42085954664 (42.09 G)
Number of raw bytes: 42101578296 (42.10 G)
Number of pcap bytes: 42949673232 (42.95 G)
Number of IPv4 packets: 53768016 (53.77 M) [99.60%]
Number of IPv6 packets: 214108 (214.11 K) [0.40%]
Number of A packets: 31005784 (31.01 M) [57.44%]
Number of B packets: 22976625 (22.98 M) [42.56%]
Number of A bytes: 14212871985 (14.21 G) [33.77%]
Number of B bytes: 27873082679 (27.87 G) [66.23%]
Average A packet load: 458.39
Average B packet load: 1213.11 (1.21 K)
=====

```

```

basicStats: Biggest Talker: 92.122.216.218: 154922 (154.92 K) [0.29%] packets
basicStats: Biggest Talker: 92.122.216.218: 224224088 (224.22 M) [0.53%] bytes
tcpFlags: Aggregated IP anomaly flags : 0x3d6e
tcpFlags: Aggregated TCP anomaly flags: 0xfe07
tcpFlags: Number of TCP scans, succ scans, retries: 175001 (175.00 K), 29259 (29.26 K), 15782
(15.78 K)
tcpFlags: Number WinSz below 1: 162725 (162.72 K) [0.38%]
tcpStates: Aggregated anomaly flags: 0xdf
icmpDecode: Number of ICMP echo request packets: 14979 (14.98 K) [12.00%]
icmpDecode: Number of ICMP echo reply packets: 2690 (2.69 K) [2.16%]
icmpDecode: ICMP echo reply / request ratio: 0.18
icmpDecode: Number of ICMPv6 echo request packets: 1440 (1.44 K) [9.63%]
icmpDecode: Number of ICMPv6 echo reply packets: 951 [6.36%]
icmpDecode: ICMPv6 echo reply / request ratio: 0.66
dnsDecode: Number of DNS packets: 237597 (237.60 K) [0.44%]
dnsDecode: Number of DNS Q packets: 125260 (125.26 K) [52.72%]
dnsDecode: Number of DNS R packets: 112337 (112.34 K) [47.28%]
dnsDecode: Aggregated status: 0xe72f
httpSniffer: Number of HTTP packets: 36614826 (36.61 M) [67.83%]
httpSniffer: Number of HTTP GET requests: 426825 (426.82 K) [1.17%]
httpSniffer: Number of HTTP POST requests: 39134 (39.13 K) [0.11%]
httpSniffer: HTTP GET/POST ratio: 10.91
httpSniffer: Aggregated status flags : 0x003f
httpSniffer: Aggregated anomaly flags: 0x5143
httpSniffer: Aggregated content flags: 0x007a
httpSniffer: Aggregated mime type : 0x80ef
connStat: Max unique number of IP source connections: 275532 (275.53 K)
connStat: Max unique number of IP destination connections: 301252 (301.25 K)
connStat: Max unique number of IP source/destination connections: 1242 (1.24 K)
connStat: Max unique number of source IP / destination port connections: 1521 (1.52 K)
connStat: prtcon/sdcon, prtcon/scon: 1.224638, 0.005520
connStat: Source IP with max connections: X.Y.Z.U: 1515 (1.51 K) connections
connStat: Destination IP with max connections: V.W.A.B: 3241 (3.24 K) connections
-----
Headers count: min: 4, max: 14, average: 7.10
Max VLAN header count: 1
Max MPLS header count: 2
Number of LLC packets: 285 [0.00%]
Number of GRE packets: 27670 (27.67 K) [0.05%]
Number of Teredo packets: 213877 (213.88 K) [0.40%]
Number of AYIYA packets: 231 [0.00%]
Number of IGMP packets: 401 [0.00%]
Number of ICMP packets: 124800 (124.80 K) [0.23%]
Number of ICMPv6 packets: 14946 (14.95 K) [0.03%]
Number of TCP packets: 43273340 (43.27 M) [80.16%]
Number of UDP packets: 10311931 (10.31 M) [19.10%]
Number of IPv4 fragmented packets: 19155 (19.16 K) [0.04%]
Number of IPv6 fragmented packets: 9950 (9.95 K) [4.65%]
~~~~~
Number of processed flows: 1439153 (1.44 M)
Number of processed A flows: 1209728 (1.21 M) [84.06%]
Number of processed B flows: 229425 (229.43 K) [15.94%]
Number of request flows: 930935 (930.93 K) [64.69%]
Number of reply flows: 508218 (508.22 K) [35.31%]
Total A/B flow asymmetry: 0.68
Total req/rply flow asymmetry: 0.29
Number of processed packets/flows: 37.51
Number of processed A packets/flows: 25.63
Number of processed B packets/flows: 100.15
Number of processed total packets/s: 237884.31 (237.88 K)

```

```

Number of processed A+B packets/s: 237884.31 (237.88 K)
Number of processed A   packets/s: 136633.20 (136.63 K)
Number of processed   B packets/s: 101251.10 (101.25 K)
~~~~~
Number of average processed flows/s: 6341.92 (6.34 K)
Average full raw bandwidth: 1484232320 b/s (1.48 Gb/s)
Average snapped bandwidth : 1483681536 b/s (1.48 Gb/s)
Average full bandwidth : 1483899904 b/s (1.48 Gb/s)
Max number of flows in memory: 262144 (262.14 K) [100.00%]
Number of flows terminated by autopilot: 816044 (816.04 K) [56.70%]
Memory usage: 2.73 GB [4.05%]
Aggregate flow status: 0x01003cfad298fb04
[WRN] L3 SnapLength < Length in IP header
[WRN] L4 header snapped
[WRN] Consecutive duplicate IP ID
[WRN] IPv4/6 fragmentation header packet missing
[WRN] IPv4/6 packet fragmentation sequence not finished
[INF] IPv4
[INF] IPv6
[INF] IPv4/6 fragmentation
[INF] VLAN encapsulation
[INF] MPLS encapsulation
[INF] L2TP encapsulation
[INF] PPP/HDLC encapsulation
[INF] GRE encapsulation
[INF] AYIYA tunnel
[INF] Teredo tunnel
[INF] CAPWAP/LWAPP tunnel
[INF] SSDP/UPnP flows
[INF] Ethernet flows
[INF] SIP/RTP flows
[INF] Authentication Header (AH)
[INF] Encapsulating Security Payload (ESP)
[INF] TOR addresses

```

Listing 1: A sample *Tranalyzer2* final report including encapsulation warning, Hash Autopilot engagement when hash table full

T2 runs in IPv4 mode, but warns the user that there is IPv6 encapsulated. Note that the new [Hash Autopilot](#) warns you when the main hash map is full. T2 then removes the oldest Flow and continues processing your pcap. To avoid that, run T2 again, but this time, use the `-f 5` option as indicated in the warning message: `[INF] Hash Autopilot: Fix: Invoke Tranalyzer with '-f 5'`
`t2 -r ~/wurst/data/knoedel.pcap -w ~/results -f 5` or just let it run to the finish.

2.14 Monitoring Modes During Runtime

If debugging is deactivated or the verbose level is zero (see Section 2.9), *Tranalyzer2* prints no status information or end report. Interrupt signal has been introduced to force intermediate status information to `stdout`. Appropriate Unix commands and their effect are listed below:

| Command | Description |
|-----------------------------|--|
| <code>kill -USR1 PID</code> | T2 sends configured monitoring report to <code>stdout</code> |
| <code>kill -USR2 PID</code> | T2 toggles between on demand and continuous monitoring operation |

The script `t2stat` has the same function as `kill -USR1 PID`. An example of a typical signal requested report

(MACHINE_REPORT=0) is shown in Listing 2.

```

@      @
|      |
=====vVv==(a      a)==vVv=====
=====\\      /=====
=====\\      /=====
                        oo
USR1 A type report: Tranalyzer 0.8.2 (Anteater), Tarantula. PID: 16355
PCAP time: 1291753261.106203 sec (Tue 07 Dec 2010 20:21:01 GMT)
PCAP duration: 35.659471 sec
Time: 1546607860.487896 sec (Fri 04 Jan 2019 14:17:40 CET)
Elapsed time: 21.399345 sec
Total bytes to process: 42949673232 (42.95 G)
Percentage completed: 16.01%
Total bytes processed so far: 6875527168 (6.88 G)
Remaining time: 112.276935 sec (1m 52s)
ETF: 1546607972.764831 sec (Fri 04 Jan 2019 14:19:32 CET)
Number of processed packets: 8576716 (8.58 M)
Number of processed bytes: 6738299436 (6.74 G)
Number of raw bytes: 6740587316 (6.74 G)
Number of IPv4 packets: 8543638 (8.54 M) [99.61%]
Number of IPv6 packets: 33035 (33.03 K) [0.39%]
Number of A packets: 4916815 (4.92 M) [57.33%]
Number of B packets: 3659901 (3.66 M) [42.67%]
Number of A bytes: 2267361054 (2.27 G) [33.65%]
Number of B bytes: 4470938382 (4.47 G) [66.35%]
Average A packet load: 461.14
Average B packet load: 1221.60 (1.22 K)
-----
tcpFlags: Number of TCP scans, succ scans, retries: 9533, 4061, 2125
icmpDecode: Number of ICMP echo request packets: 2610 (2.61 K) [11.39%]
icmpDecode: Number of ICMP echo reply packets: 703 [3.07%]
dnsDecode: Number of DNS packets: 37156 (37.16 K) [0.43%]
dnsDecode: Number of DNS Q packets: 18735 (18.73 K) [50.42%]
dnsDecode: Number of DNS R packets: 18421 (18.42 K) [49.58%]
dnsDecode: Aggregated status: 0x0201
httpSniffer: Number of HTTP packets: 5833252 (5.83 M) [68.01%]
connStat: Max unique number of IP source connections: 61971 (61.97 K)
connStat: Max unique number of IP destination connections: 67895 (67.89 K)
connStat: Max unique number of IP source/destination connections: 712
connStat: Max unique number of source IP / destination port connections: 712
connStat: IP prtcon/sdcon, prtcon/scon: 1.000000, 0.011489
-----
Headers count: min: 4, max: 13, average: 6.95
Max VLAN header count: 1
Max MPLS header count: 2
Number of LLC packets: 43 [0.00%]
Number of GRE packets: 3677 (3.68 K) [0.04%]
Number of Teredo packets: 33001 (33.00 K) [0.38%]
Number of AYIYA packets: 34 [0.00%]
Number of IGMP packets: 58 [0.00%]
Number of ICMP packets: 20491 (20.49 K) [0.24%]
Number of ICMPv6 packets: 2426 (2.43 K) [0.03%]
Number of TCP packets: 6907067 (6.91 M) [80.53%]
Number of UDP packets: 1608944 (1.61 M) [18.76%]
Number of IPv4 fragmented packets: 2511 (2.51 K) [0.03%]
Number of IPv6 fragmented packets: 540 [1.63%]
~~~~~

```

```

Number of processed flows: 270826 (270.83 K)
Number of processed A flows: 228096 (228.10 K) [84.22%]
Number of processed B flows: 42730 (42.73 K) [15.78%]
Number of request flows: 224453 (224.45 K) [82.88%]
Number of reply flows: 46373 (46.37 K) [17.12%]
Total A/B flow asymmetry: 0.68
Total req/rply flow asymmetry: 0.66
Number of processed packets/flows: 31.67
Number of processed A packets/flows: 21.56
Number of processed B packets/flows: 85.65
Number of processed total packets/s: 240517.21 (240.52 K)
Number of processed A+B packets/s: 240517.21 (240.52 K)
Number of processed A packets/s: 137882.45 (137.88 K)
Number of processed B packets/s: 102634.76 (102.63 K)
~~~~~
Number of average processed flows/s: 7594.78 (7.59 K)
Average full raw bandwidth: 1512212608 b/s (1.51 Gb/s)
Average snapped bandwidth : 1511699328 b/s (1.51 Gb/s)
Average full bandwidth : 1511878912 b/s (1.51 Gb/s)
Fill size of main hash map: 236513 [90.22%]
Max number of flows in memory: 236517 (236.52 K) [90.22%]
Memory usage: 2.49 GB [3.69%]
Aggregate flow status: 0x01003872d298fb04
[WRN] L3 SnapLength < Length in IP header
[WRN] Consecutive duplicate IP ID
[WRN] IPv4/6 fragmentation header packet missing
[INF] IPv4
[INF] IPv6
[INF] IPv4/6 fragmentation
[INF] VLAN encapsulation
[INF] MPLS encapsulation
[INF] L2TP encapsulation
[INF] PPP/HDLC encapsulation
[INF] GRE encapsulation
[INF] AYIYA tunnel
[INF] Teredo tunnel
[INF] CAPWAP/LWAPP tunnel
[INF] SSDP/UPnP flows
[INF] Ethernet flows
[INF] SIP/RTP flows
[INF] Authentication Header (AH)
[INF] Encapsulating Security Payload (ESP)
[INF] TOR addresses
=====

```

Listing 2: A sample Tranalyzer2 human readable report aggregate mode

Listing 3 illustrates the output of the header line and subsequent data lines generated when MACHINE_REPORT=1.

| %RepTyp | Time | Dur | memUsq[KB] | fillSzHashMap | Flws | AFlws | BFlws | SIP | DIP | SDIP | Prts |
|------------|-------------------|------------|------------|---------------|----------|----------|-----------|-----------|-----------|------------|------------|
| Pkts | APkts | BPkts | V4Pkts | V6Pkts | VxPkts | Byts | AByts | BByts | ARPPkts | RARPPkts | |
| ICMPPkts | EchoReq | EchoRep | DnsPkts | DnsQPkts | DnsRPkts | HttpPkts | FrgV4Pkts | FrgV6Pkts | | | |
| Alrms | TCPScn | TCPSScn | TcpRtry | RawBndWdth | Fave | GlblWrn | ICMP | IGMP | TCP | UDP | SCTP |
| 0x0042 | 0x00fe | 0x0806 | 0x0800 | 0x86dd | | | | | | | |
| USR1MR_A | 1175364020.458606 | | 103.877184 | 1856384 | | 6456055 | 518079 | 360556 | 157523 | 92140 | |
| 153825 | 765 | 1706 | 9453374 | 5485249 | 3967298 | 9452802 | 0 | 0 | 680258574 | 2271454049 | 3288136123 |
| 0 | 0 | 492113898 | 1244 | 0 | 0 | 3309 | 20599 | 0 | 0 | 26244 | 22467 |
| 428175.000 | 2.230065 | 0x02030052 | 49211 | 102 | 8186447 | 1169051 | 0 | 284 | 0 | 0 | 9452802 |
| USR1MR_A | 1175364043.564798 | | 126.983376 | 2030752 | | 6945643 | 680946 | 462230 | 218716 | 118181 | |
| 183651 | 1164 | 3626 | 13104550 | 7235588 | 5867866 | 13103799 | 0 | 0 | 944828393 | 3007472308 | |
| 4847491640 | 0 | 066918 | 5110 | 1695 | 0 | 0 | 0 | 4738 | 28844 | 0 | 0 |
| | | | | | | | | | 33600 | 32326 | |

```

23344 494874.688 3.115120 0x02030052 66918 130 11364263 1611189 0 372 0 0
13103799 0
USR1MR_A 1175364072.842855 156.261433 2184732 7103403 860105 573862 286243 148357
217272 1164 3626 17716953 9443884 8271657 17715947 0 0 1278800701 3927484585
6829832070 0 086334 6700 2226 0 0 0 6572 38940 0 0 48853 45035
34063 550743.562 3.115120 0x02030052 86334 177 15394866 2158059 0 492 0 0
17715947 0

```

Listing 3: A sample Tranalyzer2 machine report aggregate mode

2.14.1 Configuration for Monitoring Mode

To enable monitoring mode, configure Tranalyzer as follows:

| main.h | | tranalyzer.h | |
|---------|---------------|--------------|------------------|
| #define | MONINTPSYNC 1 | #define | VERBOSE 0 |
| #define | MONINTTMPCP 1 | #define | DIFF_REPORT 1 |
| #define | MONINTTHRD 1 | #define | MACHINE_REPORT 1 |

The following plugins contribute to the output:

- basicStats
- connStat
- dnsBCmp
- dnsDecode
- httpBCmp
- icmpDecode
- protoStats

The generated output is illustrated in Figure 3. The columns are as follows:

| | | |
|------------------|--------------|----------------|
| 1. RepType | 14. APkts | 27. DnsPkts |
| 2. Time | 15. BPkts | 28. DnsQPkts |
| 3. Dur | 16. V4Pkts | 29. DnsRPkts |
| 4. memUsage[KB] | 17. V6Pkts | 30. HttpPkts |
| 5. fillSzHashMap | 18. VxPkts | 31. FrgV4Pkts |
| 6. Flows | 19. Byts | 32. FrgV6Pkts |
| 7. AFlows | 20. AByts | 33. Alrms |
| 8. BFlows | 21. BByts | 34. TCPScn |
| 9. SIP | 22. ARPPkts | 35. TCPSScn |
| 10. DIP | 23. RARPPkts | 36. TcpRtry |
| 11. SDIP | 24. ICMPPkts | 37. RawBndWdth |
| 12. Prts | 25. EchoReq | 38. Fave |
| 13. Pkts | 26. EchoRep | 39. GlblWrn |

When capturing from a live interface, the following three columns are output (between `Dur` and `memUsage[KB]`):

- PktsRec
- PktsDrp
- IfDrp

2.14.2 Monitoring Mode to syslog

In order to send monitoring info to a syslog server T2 must be configured in machine mode as indicated above. Then the output has to be piped into the following script:

```
t2 -D ... -w ... | awk -F"\t" '{ print "<25> ", strftime("%b %d %T"), "Monitoring: " $0; }' | \
nc -u w.x.y.z 514
```

Netcat will send it to the syslog server at address `w.x.y.z`. Specific columns from the monitoring output can be selected in the awk script.

2.14.3 RRD Graphing of Monitoring Output

The monitoring output can be stored in a RRD database using the `rrdmonitor` script. To start creating a RRD database, launch Tranalyzer2 (in monitoring mode) as follows:

```
./tranalyzer2/src/tranalyzer -r file.pcap | ./scripts/rrdmonitor
```

Or for monitoring from a live interface:

```
sudo ./tranalyzer2/src/tranalyzer -i eth0 | ./scripts/rrdmonitor
```

Plots for the various fields can then be generated using the `rrdplot` script. To specify intervals, use `s` (seconds), `m` (minutes), `h` (hour), `d` (day), `w` (week), `mo` (month), `y` (year). For example, to plot the data from the last two weeks, use `-i 2w` or `-s -2w`.

An example graph is depicted in Figure 3.




Figure 3: T2 monitoring using RRD

2.15 Cancellation of the Sniffing Process

Processing of a pcap file stops upon end of file. In case of live capture from an interface Tranalyzer2 stops upon `CTRL+C` interrupt or a `kill -9 PID` signal. The disconnection of the interface cable will stop Tranalyzer2 also after a timeout of 182 seconds. The console based `CTRL+C` interrupt does not immediately terminate the program to avoid corrupted entries in the output files. It stops creating additional flows and finishes only currently active flows. Note that waiting the termination of active flow depends on the activity or the lifetime of a connection and can take a very long time. In order to mitigate that problem the user can issue the `CTRL+C` for `GI_TERM_THRESHOLD` times to immediately terminate the program.